

how-to-use-[ai.com](https://how-to-use-ai.com) |

Book 1 of 5

AI for Normal People

Understanding Artificial Intelligence Without the Hype

No technical
skills required



Danijel-James Wynyard-McClay

How-To-Use-AI.com

BOOK 1

AI for Normal People

PREVIEW EDITION

This free preview contains Chapters 1–3 of the book.

The text may change before final publication.

Danijel-James Wynyard-McClay

A free preview of a book in progress.

AI for Normal People

Understanding Artificial Intelligence Without the Hype

Danijel-James Wynyard-McClay

how-to-use-[ai](https://how-to-use-ai.com).com

RePass Cloud Pty Ltd

AI for Normal People: Understanding Artificial Intelligence Without the Hype

by Danijel-James Wynyard-McClay

Copyright © 2026 Danijel-James Wynyard-McClay. All rights reserved.

Published in Australia by RePass Cloud Pty Ltd. 3 Pioneer Street, Findon SA 5023, Australia.

how-to-use-ai.com

Paperback **ISBN 978-1-7649948-0-4**

EPUB **ISBN 978-1-7649948-2-8**

PDF **ISBN 978-1-7649948-1-1**

First edition, October 2026

AI products change quickly. Product names, features and behaviour described in this book were checked at the time of writing and may have changed since. This book gives general information, not professional, legal or financial advice.

Product names mentioned in this book are trademarks of their respective owners. This book is independent and is not endorsed by them.

No part of this publication may be reproduced, stored or transmitted in any form without the prior written permission of the publisher, except as permitted under the Copyright Act 1968 (Cth).

Preview edition. The text may change before final publication.

For Johnny,

*who has endured more midnight conversations about AI, technology, and whatever else
my ADHD brain has decided is fascinating that night than any husband reasonably
should.*

*You listen while I excitedly explain one more idea when we should both be asleep,
knowing full well that “one more thing” is never actually one more thing.*

And somehow, without fail, you still love me every day.

This book is for you.

**Any sufficiently advanced technology is
indistinguishable from magic.**

— Arthur C. Clarke

Contents

How This Book Works	vi
1 You've Already Been Using AI	1
2 Why Everyone Suddenly Talks About AI	11
3 What AI Can Actually Do	26
Notes	43
Further Reading	46

How This Book Works

Most of this book is ordinary reading: plain-English explanations, everyday examples, and the occasional table. Every so often, though, you'll come across a shaded or outlined box with a small coloured label in its top corner. Those boxes pull out the handful of things worth stopping for.

There are five kinds, and each one always means the same thing.

Key Idea

The one thing to remember from a section. If you only skim a chapter, read these.

Try This

A small, practical thing you can do right now, usually in a few minutes and without any technical setup. You don't have to do them all, but they're the quickest way to turn reading into experience.

Watch Out

A risk, a limitation, or a common misunderstanding. These are the places where AI tools most often catch people out, so they're worth slowing down for.

Recap

A short summary at the end of each chapter, often with a question to think about. Use it to check that the chapter landed before you move on.

[Author reflection: These boxes are where I step out of the explanation and share what something has looked like in practice through my own work with AI, software and technology. The experiences are genuine rather than invented to fit the chapter, and appear only where I have something useful to add from first-hand experience.]

You can read the chapters in order, which is how the ideas build, or dip into the chapter that matches what you need today. Either way, the boxes are there to help you find your footing quickly.

You've Already Been Using AI

"Most people think they have never used AI. In reality, most people use it dozens — sometimes hundreds — of times every day."

Artificial intelligence sounds intimidating because most of us picture robots, supercomputers, or something out of a science fiction film. The term itself feels technical, expensive, and far removed from ordinary life. But here's the thing: you've probably used AI today already, possibly several times before you finished breakfast.

If you've watched a recommendation on Netflix, followed directions in Google Maps, ignored a piece of spam email, scrolled through TikTok, spoken to Siri or Alexa, received a fraud alert from your bank, or bought something an online store suggested for you — you've used AI. Not an experimental version, and not some "future" AI still in the lab. Real AI systems, already running at massive scale, quietly doing their job in the background of things you use every day.

The goal of this chapter is to remove the mystery around AI before we go any further, because the truth is a lot simpler than the hype suggests. At its core, AI is software that finds patterns, predicts likely outcomes, and improves its decisions as it sees more data. That's the foundation underneath nearly every AI system you'll interact with, from the mundane to the genuinely impressive.

Key Idea

AI is not magic. It is pattern recognition at scale.

AI Is Already Everywhere

Most AI systems are invisible. They don't introduce themselves, and they certainly don't announce that "Artificial Intelligence" is operating in the background. Instead, they quietly filter, rank, predict, recommend, flag anomalies, interpret language, and recognise images or sounds — all without asking for credit.

The reason so many people suddenly became aware of AI in the last few years isn't that it appeared out of nowhere. It's that conversational AI tools made it visible and interactive for the first time. The underlying technology, though, has been running inside consumer products for a long time. The examples below are the ones most people recognise instantly once they're pointed out.

Streaming Recommendations

When a streaming platform recommends a movie or series, it's trying to answer one question: what is this person most likely to watch next? To do that, it studies what you watched before, what similar users watched, how long you stuck with something, what you skipped, what time of day you tend to watch, and which genres you keep returning to. It then compares your behaviour against millions of other viewing patterns to find the closest match.

This isn't the system "understanding" entertainment the way a person does — it's statistical prediction. Underneath, it's really just saying: people with similar viewing behaviour to yours often choose this next.

Navigation and Traffic Prediction

When Google Maps suggests a faster route, it's drawing on an enormous amount of live and historical information: road speeds, accidents, traffic density, construction, time of day, typical commuting behaviour, and even nearby events that might cause congestion. From all of that, it predicts which route will most likely get you there fastest.

The app isn't "thinking" like a human driver weighing up shortcuts. It's processing patterns across a dataset far larger and faster-changing than any person could track — which is exactly why it can reroute you around a traffic jam three suburbs away before you'd have any way of knowing it was there.




Figure 1-1: Recommendation engine feedback loop: watching content produces behavioural data, which the AI uses to predict interests and display recommendations; later interactions feed back into the system.

Spam Filters

Email spam filtering is one of the oldest large-scale consumer AI systems still running today, quietly processing billions of emails every single day. It looks for patterns like suspicious wording, unusual links, sending behaviour, known scam structures, malicious attachments, and abnormal message volume, then predicts how likely a given message is to be unwanted or dangerous.

Most people never think about spam filtering anymore, and that's precisely the point — some of the most successful AI systems become invisible once they're reliable enough. You only notice a spam filter when it gets something wrong.

Social Media Algorithms

Social media platforms lean on AI heavily, because attention is the product they're selling to advertisers. Apps like TikTok, YouTube, Instagram, and Facebook are constantly predicting what you'll click, what you'll watch longer, what you'll share, what you'll react to emotionally, and what will keep you engaged a little longer. The recommendation system adapts continuously based on your behaviour — even a few extra seconds spent watching one video becomes data it learns from.

This is part of why social media can feel unusually “accurate” at times, almost like it's reading your mind. It isn't. It's observing patterns at a scale no individual human could track, and reacting to them faster than you'd notice.

Watch Out

Accuracy doesn't mean understanding. A system can get extremely good at predicting your behaviour without ever understanding meaning, emotion, truth, or context the way a person does. It's matching patterns, not reading your mind.

Voice Assistants

When you speak to Siri or Alexa, several AI systems are working together behind that single sentence. Say “set a timer for 15 minutes,” and the assistant has to convert your speech into text, work out what you actually want, pull the number 15 out of the sentence, recognise “minutes” as the unit, and then execute the right action — all in about a second. That's speech recognition, language interpretation, intent prediction, response generation, and action execution, chained together so smoothly it feels like one simple step.

Again, none of that is magic. It's highly refined pattern recognition, layered several systems deep.

Banking Fraud Detection

Banks rely on AI because modern financial systems generate far more transactions than any team of humans could review by hand. Their systems watch for unusual spending locations, unexpected purchase sizes, abnormal login behaviour, impossible travel scenarios, and rapid bursts of transactions.

If your card is used in Sydney and then, twenty minutes later, in another country, the system doesn't need a human to tell it something's wrong — it flags the pattern immediately, because it's really just asking whether this behaviour matches your normal pattern. That's anomaly detection: identifying behaviour that differs from what's expected, at a speed that makes it useful for actually stopping fraud rather than just reporting it afterward.

Online Shopping Recommendations

Online stores lean on AI heavily because recommendations directly affect what you buy. When a site tells you “customers also bought...”, it's drawing on purchase combinations, your browsing history, abandoned carts, similarities to other shoppers, seasonal patterns, and general product popularity. Once again, this is prediction, not certainty and not human understanding — just likelihood, calculated from historical patterns across millions of other shoppers who looked a lot like you.

What AI Actually Is

At its simplest, AI is software designed to identify patterns and make predictions using data. That definition alone removes a lot of the mystery: AI systems don't need consciousness, emotion, self-awareness, or independent thought to be useful, or even commercially transformative.

A Specialised Tool, Not a General Mind

Most of the AI you actually encounter is a specialised tool, trained to do one narrow task extremely well, rather than a single all-purpose intelligence. The spam filter that protects your inbox has no idea how to suggest a route home, and the navigation app rerouting you around traffic has no idea what's in your inbox at all. Each one is narrow on purpose — a specialist, not a generalist.

Task	AI Purpose
Spam filtering	Predict unwanted email
Navigation	Predict fastest route
Recommendations	Predict user interest
Fraud detection	Predict suspicious behaviour
Voice assistants	Predict user intent




Figure 1-2: Fraud detection compares new card activity with your normal pattern: purchases in Sydney, Parramatta and Chatswood fit the pattern and are approved, while a purchase in Singapore twenty minutes after one in Sydney is impossible travel and is flagged as a fraud risk.

Pattern Recognition, at Scale

Pattern recognition — the process of spotting recurring structures, behaviours, or relationships in data — is something humans already do naturally. You recognise a friend's handwriting, notice a favourite café is busier on Fridays, or sense when a conversation is about to turn awkward, all through the same basic skill: spotting a pattern you've seen before. AI systems do the same thing mathematically, at a scale no person could manage — comparing millions of data points instead of a lifetime of personal experience — which is exactly why this technology has quietly worked its way into so much of modern life without ever needing to look like science fiction.

Why AI Feels Suddenly New

AI didn't appear overnight. What changed was accessibility.

Quiet for Years

For years, most AI systems operated silently inside large platforms and enterprises, doing their pattern-matching work behind the scenes with no interface a regular person would ever see. A bank's fraud system, an airline's pricing engine, a logistics company's delivery-route planner — all running AI, all invisible to the public, all without ever needing to explain themselves to an ordinary customer.

Visible Almost Overnight

Modern generative AI tools changed that, because for the first time, people could interact with AI directly through conversation — asking questions, generating images, writing documents, summarising information, brainstorming ideas — instead of only benefiting from predictions made invisibly on their behalf. That's what made AI visible to a huge number of consumers all at once. But the underlying concept barely changed: it's still pattern recognition, prediction, and probability. What's new is the interface, not the idea underneath it.

Author Reflection

I don't remember exactly when Spotify's recommendations stopped feeling like recommendations and started feeling slightly intrusive.

But I remember noticing the change.

It wasn't when it suggested another artist from a genre I already listened to. That's hardly witchcraft. If I've spent the afternoon listening to several rock bands and Spotify recommends another rock band, nobody needs to call Skynet.

It was when it started making jumps I would not have made myself.

One example that stuck with me was Porter Robinson's "Goodbye To A World". On paper, it was not an obvious recommendation based on what I thought I listened to. Different artist, different corner of music, not something I would have searched for myself.

And yet within about thirty seconds I understood why it had put it in front of me.

There was something familiar in it. Maybe the vocal style, the production, the rhythm, the atmosphere, the way the chorus built — some combination of characteristics that connected it to music I already liked even though I had never consciously made that connection myself.

That was the bit that got my attention.

Spotify had apparently worked out something about my taste that I had never actually told it.

At least, not deliberately.

I had told it every time I let a song finish instead of skipping it. Every time I replayed something. Every artist I searched for. Every track I saved. Every playlist I returned to. Every song I abandoned twelve seconds in because apparently twelve seconds was enough to conduct a complete artistic review.

Individually, none of those actions meant very much.

Collectively, I had spent years feeding it an absurdly detailed description of what I liked without ever filling out a questionnaire.

And it wasn't simply learning, "DJ likes this artist."

That would be easy. And if you know me - that wouldn't work. I might like this song by that artist - but I have no idea what their other songs are or I like this song and no clue who the artist even was or ever bothered to glance at the screen to find out!

It was learning the patterns underneath that choice: people who repeatedly listen to these things, skip those things, replay these particular tracks and wander into these other genres often respond well to something over here.

That is considerably more interesting.

It also explains why changing music services can feel strangely like starting again.

You can move your playlists. You can rebuild your library. The songs themselves are mostly available everywhere.

What doesn't come with you quite so neatly is the history.

Years of skips, repeats, searches, discoveries, temporary obsessions and songs you played forty times before inexplicably deciding never to acknowledge their existence again.

That accumulated behaviour has value because it lets the recommendation system build a model that becomes increasingly specific to you.

Nobody ever sat me down and said, “We are going to train an AI system to predict your musical taste.”

I just listened to music.

That is what makes it such a useful example of how AI existed in ordinary life long before most people started talking about AI.

There was no chatbot. No prompt box. No glowing robot eye announcing that machine learning had entered the building.

Just Spotify quietly putting a song in front of me that I would never have looked for myself.

Then, somewhat annoyingly, being absolutely right.

The Problem With AI Hype

Part of what makes AI discussions confusing is that people naturally exaggerate new technology.

Common Exaggerations

It’s easy to slip into describing AI as though it understands everything, thinks independently, replaces human judgment entirely, operates like a digital brain, or has some kind of intent or awareness. Most current AI systems do none of that. They’re advanced prediction systems trained on enormous amounts of data — some of them incredibly powerful, some genuinely transformative for the businesses and people using them, but all still fundamentally pattern-based.

Why the Distinction Matters

That distinction matters more than it might seem, because understanding what AI actually is makes it far easier to evaluate claims realistically, avoid fear-based narratives, understand where the real limitations sit, use AI effectively, and recognise where human judgment still needs to be in the loop.

Where the Confusion Comes From

Part of the confusion comes from language. It's convenient to say an AI system “thinks,” “understands,” “learns,” or “knows” something — but those are human words borrowed for a process that's really mathematical pattern matching. They're useful shorthand, but taken literally, they set up expectations the technology can't actually meet.

Watch Out

A few myths are worth retiring early, before they shape how you read the rest of this book. AI isn't only robots and science fiction — it's already inside everyday software, and most of the time it works quietly in the background where you never see it. It doesn't think like a human; most AI systems predict patterns from data. It isn't magic; it's mathematics, data, and software engineering. And it isn't always correct. It predicts likely outcomes, and a prediction can be wrong.

Try This

For one day, notice where AI might be operating in the background of things you do — recommended videos, map route suggestions, email filtering, banking alerts, shopping recommendations, autocomplete, search rankings, voice assistant responses. The goal isn't to get technical. It's to notice that AI is already ordinary.

Core Takeaway

AI is not magic, not science fiction, and not something that suddenly appeared out of nowhere. Most people have already been using it every day for years, often without a second thought. At its core, AI is pattern recognition at scale — and that single idea is the foundation for everything else in this book series.

Recap

Familiar tools already use AI-style pattern prediction — streaming recommendations, navigation, spam filters, social feeds, voice assistants, fraud detection, and shopping suggestions all work the same basic way: find patterns, predict likely outcomes, improve with more data. AI is not magic. It is pattern recognition at scale. Where have you already used it without calling it AI?

Why Everyone Suddenly Talks About AI

If you finished Chapter 1 and thought, “okay, so AI has been quietly running my spam filter and my banking app for years — so why does it suddenly feel like everyone is talking about it for the first time?” — that’s exactly the right question to ask.

You’re not wrong. AI did not suddenly appear. The most visible change was that ordinary people could open a text box, type a plain-English question, and get a useful-looking answer back. That interface mattered enormously because it turned AI from something most people experienced indirectly into something they could try for themselves.

But “the interface changed” is only the beginning of the honest answer. The systems behind that interface had also become more capable. Companies had gained access to larger collections of data, far more computing power, and years of research advances. Investors were willing to spend extraordinary sums building the technology and placing it inside products people already used. News coverage, workplace anxiety, and competitive pressure then amplified the effect.

So the current AI wave has several causes, not one. Conversation opened the front door. Better models made the visit impressive. Money put AI into almost every shop window. Human curiosity, hope, fear, and marketing did the rest.

Key Idea

The interface changed from buttons and menus to conversation, but the sudden attention came from capability, access, investment, and publicity arriving together.

From Quiet Prediction to Loud Conversation

Think back to the examples from Chapter 1. Your spam filter predicts which emails are unwanted. Your streaming app predicts what you'll watch next. Your bank predicts which transactions look suspicious. In every one of those cases, the AI is working *for* you, in the background, without you ever typing a request to it directly. You didn't ask your bank's fraud system a question — it just watched patterns and quietly flagged anything unusual.

Now compare that with typing “explain this bill to me in simple terms” into a chat window and watching a written answer appear a few seconds later. Pattern prediction still sits underneath the result, but the scale and flexibility are different. The output can be a paragraph shaped around your request rather than a single flag, score, route, or ranked list. AI did not stop doing quiet prediction; that kind of AI still surrounds us. A more open-ended form of prediction simply became visible beside it.

What Changed: Capability, Access, and Interface

It's worth being precise about what actually changed, because two opposite stories are commonly told and both are misleading. In one story, AI suddenly became intelligent in 2022. In the other, nothing important changed except the design of a website. The truth sits between them. Language technology had been improving for years, and the models available by 2022 were meaningfully more capable than earlier systems. What changed suddenly was public access: a complicated research and engineering stack appeared behind a remarkably simple chat box.

Previously, direct experimentation with leading language systems usually required technical skill, specialist access, or a product built by somebody else. After the public release of ChatGPT on 30 November 2022, millions of people could encounter a conversational language model through an ordinary website.¹ Access was not literally universal—cost, connectivity, language, disability, geography, and account restrictions still matter—but the barrier was dramatically lower.

Filling Out a Form vs. Asking a Question

It helps to see the shift side by side. Imagine you want to return a pair of shoes you bought online because they don't fit. The traditional way: log in, find “Orders,” find the right order among several, click “Return,” select a reason from a dropdown that doesn't quite match your situation, print a label, and wait for a confirmation email. Every step exists because the software needs your intent broken down into fields it already knows how to handle.

The conversational way: type “I bought these shoes last week and they don't fit; how do I send them back?” into a chat window. A well-designed assistant might explain the policy, ask a clarifying question, or direct you to the correct step. If it is connected to the retailer's order system and has permission, it might even prepare the return. If it is only a general chatbot, however, it cannot know your particular order or guarantee the retailer's rules. The conversation feels natural, but the system's access and authority still determine what it can actually do.




Figure 2-1: A traditional interface makes you translate what you want into menus, buttons, forms and dropdowns, while a conversational interface lets you type a question and get a written answer back.

That small shift — from “find the button that means what I mean” to “just say what you mean” — is a big part of why generative AI felt so immediately usable to people who’d never think of themselves as technical. It removed a translation step most of us had stopped noticing we were doing.

Try This

Next time you have a routine question — how to change a booking, what a bill’s line item means, how to phrase an awkward email — try asking a conversational AI tool in plain English before you go looking for the right menu, form, or search result. Notice how much of the “work” was just figuring out where to put your question.

Plain English: What a Language Model Actually Does

A **language model** learns statistical relationships in language and uses the text already present to estimate plausible continuations. Modern systems usually work with small units called *tokens*, which may be whole words, pieces of words, or punctuation. They generate a response token by token rather than retrieving a guaranteed finished answer from a shelf.²

Phone autocomplete is a useful first comparison, but it should not be taken too literally. A large language model has many more learned relationships, a much larger context, and training designed to make it respond to instructions. Nor should we say it has “read the internet” as though it studied, understood, and remembered every page the way a person might. Training data varies by model, and a model’s internal patterns are not a dependable catalogue of its sources.

Generative AI vs. Predictive AI

Chapter 1’s examples—recommendations, navigation, spam filters, and fraud detection—usually produce a constrained result: a ranking, route, category, score, or warning. **Generative AI** produces new arrangements of content such as text, images, audio, video, or code. When you ask a chat tool to draft an email or summarise a document, it is not normally fetching one complete matching answer from a database. It is generating an output from learned patterns and the context you supplied.

This is a useful distinction, not a perfect border. Generative tools still make predictions, and many products combine generative models with search, databases, recommendation systems, and ordinary software. “Generative” describes what the model produces; it does not mean the result is original in the human sense, factually correct, or free from material found in its training data.

This combination—more capable generation, plain-language instructions, and low-friction public access—is the technical and product shift underneath the sudden explosion. Chapter 3 picks this up properly and examines what that generating ability is actually good for.

AI Has Been Loud Before

Here's something worth sitting with before going any further: this is not the first time the public got excited about AI, and it's not the first time that excitement crashed back to earth. Knowing that history doesn't tell you what happens next, but it does explain why some people who've been watching this field for decades react to the current wave with more caution than headlines might suggest is warranted.

The First Boom and Bust

The name *artificial intelligence* was proposed for a 1956 summer research project at Dartmouth College. The original proposal was strikingly ambitious: its organisers suggested that a carefully selected group could make significant progress on language, abstraction, problem-solving, and machine learning during one summer.³ That confidence helped establish a field, but early demonstrations worked inside much smaller worlds than the messy situations people expected them to master.

By the early 1970s, enthusiasm had weakened. Computers were limited, many problems were harder than predicted, and promised breakthroughs had not arrived. In Britain, the 1973 Lighthill report expressed deep pessimism about near-term progress and contributed to reduced support for AI research. That period later became known as an **AI winter**: not a time when all work stopped, but a time when confidence and funding cooled sharply.⁴

The Second Boom and Bust

AI did not stay frozen forever. Through the 1980s, **expert systems**—software that applied hand-written rules representing specialist knowledge—became a commercial success story. Some performed valuable work in narrow domains. Their limits also became expensive: knowledge was difficult to capture, exceptions multiplied, maintenance was laborious, and unusual inputs could produce bizarre answers. The United States Defense Advanced Research Projects Agency describes the resulting pattern plainly: early successes led to a boom, then hype gave way to disappointment in the late 1980s.⁵

The familiar story calls this a second AI winter. That label is useful, but history is tidier in hindsight than it was at the time. Funding and fashion shifted while important work in probabilistic methods, machine learning, robotics, and other areas continued. “Winter” describes the collapse of expectations around a dominant story; it does not mean intelligence research went into hibernation everywhere.

Deep Learning's Quiet Return

The recovery was slower and less visible to the general public. Researchers kept improving statistical machine learning and neural networks while computing hardware and available training data expanded. In 2012, a deep neural network produced a major advance in large-scale image classification. In 2017, researchers introduced the **transformer**, an architecture that processed relationships in sequences efficiently and became foundational to modern language models.⁶

Neither paper created today's chatbots by itself. Research progress rarely follows a single heroic moment. Algorithms, data, specialised chips, software tools, engineering methods, investment, and product design accumulated over years. By the time a conversational version reached the public in 2022, the result looked sudden mainly because the long build-up had happened outside most people's view.

It hadn't come from nowhere. It came from roughly seventy years of boom, bust, and quiet recovery — which is worth remembering both when the coverage turns euphoric and when it turns fearful.

Why 2022 Felt Like a Turning Point

Late 2022 is widely treated as the moment AI went from “a thing that exists” to “a thing everyone is talking about.” ChatGPT was not the first chatbot, the first language model, or the invention of generative AI. Its importance was as a public demonstration: a capable model, tuned for dialogue, offered through an interface almost anybody already understood. OpenAI itself launched it as a free research preview and explicitly warned that it could produce plausible-sounding but incorrect or nonsensical answers.⁷

That distinction matters for the rest of this book. When something feels brand new, it's worth asking: is the underlying idea actually new, or did the *access* to it just change? Very often, with AI, it's the second one.

Why This Wave Feels Different From Earlier Tech Shifts

It is fair to ask why this felt different from the arrival of smartphones or the public internet—two earlier waves also described as changing everything. We should not rewrite those technologies as mere delivery pipes. The internet changed who could publish and collaborate; smartphones changed how people created, navigated, traded, and communicated. They created new behaviour as well as faster access.

Generative AI nevertheless changes the starting point for many information tasks. Instead of finding a page and adapting it yourself, you can request a draft, explanation, image, plan, or piece of code shaped around your instructions. The answer may be wrong and the work may still require judgement, but the first-pass production happens inside the interaction. That is a genuinely different experience from browsing a list of links.

It also spreads differently. A general-purpose chat interface can appear inside search engines, office software, phones, customer-service systems, creative tools, and coding environments. The same underlying capability can therefore reach many kinds of work at once. That breadth—not a claim that the internet or smartphone did nothing comparable—helps explain the intensity of the reaction.

Watch Out

A fluent answer can make a system appear more knowledgeable and dependable than it is. NIST calls confidently presented false content *confabulation* and treats it as a natural risk of generative models, not a rare cosmetic defect.⁸ The models did improve, but no release date turned plausible language into guaranteed truth. Check important claims rather than trusting the tone.

The Money Behind the Moment

Public excitement is one half of why AI suddenly feels inescapable. The other half is money, and there's a lot of it moving around. Understanding roughly where it's going — and why — helps explain everything from the AI features suddenly bolted onto software you already use, to the breathless headlines, to some of the more valid worries covered later in this chapter.

Why Investors Rushed In

Once a conversational AI product attracted mass attention, investors did not need proof that every promised business model would work. They needed to believe that a few winners could become enormously valuable. Venture capital is built around that kind of uneven outcome: many bets can fail if a small number become dominant.

The scale is not imaginary. Stanford's 2025 AI Index estimated that private investment in generative AI reached \$33.9 billion in 2024, more than eight times its 2022 level. The same report estimated total corporate AI investment—including private investment, mergers and acquisitions, and public offerings—at \$252.3 billion for 2024.⁹ Those categories are not interchangeable, which is one reason impressive-looking totals from different articles often disagree.

Investment does prove that powerful organisations are making large bets. It does not prove that the products will become profitable, that customers will keep paying, or that predicted social changes will occur on schedule. Money can accelerate research, buy computing infrastructure, subsidise low prices, fund publicity, and place AI features in familiar software. It can make a technology difficult to ignore before the market has decided what it is worth.

The Deals That Fund Each Other

One pattern worth understanding is the strategic partnership between an AI developer and a company selling the computing infrastructure it needs. Microsoft announced a \$1 billion investment in OpenAI in 2019 alongside an exclusive Azure computing partnership. In 2023 it described Azure as OpenAI's exclusive cloud provider. Amazon later completed a \$4 billion investment in Anthropic while describing AWS as Anthropic's primary cloud provider and Amazon's Trainium and Inferentia chips as part of the arrangement.¹⁰

Commentators sometimes call arrangements like these **circular deals** because the investor is also a major supplier and commercial partner. The phrase can be useful, but it can also smuggle in a conclusion the public evidence does not establish. A disclosed partnership is not automatically fraudulent, fake revenue, or money literally travelling in a closed loop. It does mean the parties' incentives are connected.

That distinction changes how you read a headline. An infrastructure provider investing in a model company is evidence of commitment, access, and strategic alignment. It is not the same as an unrelated customer choosing a product at full price, nor is it independent proof that the broader market wants the product. Readers should ask what kind of money is being reported: equity investment, infrastructure credits, contracted spending, revenue, or profit. A large number can be real and still mean less than the headline implies.

How It Bled Into Everyday Life

None of this stayed confined to business pages and investor calls. Chatbots became homework tools, writing partners, jokes, workplace experiments, and sources of controversy. Image generators raised visible questions about art and copyright. Voice cloning made impersonation feel personal rather than theoretical. Search engines and office suites began presenting AI features to people who had never deliberately sought out an AI product.

There is evidence of rapid adoption, but it is more uneven than the cultural noise suggests. In a Pew Research Center survey conducted in early 2025, 34 per cent of United States adults said they had used ChatGPT, up from 18 per cent in 2023. Usage differed sharply by age: 58 per cent of adults under 30 reported having used it, compared with 10 per cent of those aged 65 and older.¹¹ That is fast growth. It is not "everyone," and it is not a global measure.

This is an important media-literacy lesson. A tool can dominate conversation while most people have not used it. People who work in technology, media, education, consulting, and other information-heavy fields can mistake their own environment for the whole population. Mainstream attention and mainstream habitual use are related, but they are not the same fact.

Why Businesses Are Both Excited and Anxious

That same mainstream attention lands differently once you're not just a curious individual but a business owner, manager, or employee — because once a technology becomes something ordinary people can access directly, businesses tend to react in two directions at once.

The Efficiency Case

There's excitement because a tool that can draft, summarise, translate, brainstorm, and explain on demand creates real opportunities to reduce effort. Controlled studies have found meaningful improvements on some writing, customer-support, coding, and consulting tasks. But an OECD review of the evidence stresses that results depend on the task and the user's experience, and that long-term business effects remain uncertain.¹²

That qualification matters. A support worker may draft a routine reply faster. A manager may turn notes into a first-pass summary. A small business owner may create a serviceable product description without hiring a copywriter for that draft. These are useful changes. They do not prove that every worker becomes proportionally more productive, that saved minutes become profit, or that an entire occupation can be removed without losing judgement, relationships, and accountability.

The Anxiety Case

And there's anxiety, for several overlapping reasons. Competitors might adopt it faster, so nobody wants to be left behind—even when nobody can define what “ahead” means. Employees may use public tools without permission, creating confidentiality and quality risks. Vendors may add AI features faster than organisations can evaluate them. People whose work includes drafting, summarising, or first-pass research may reasonably wonder what happens to their roles.

Here again, the evidence corrects the loudest story. The United States Census Bureau estimated that only 5.4 per cent of businesses were using AI to produce goods or services by February 2024, although adoption varied by industry and business size. Its survey also found that few AI-using firms reported employment reductions at that stage.¹³ The definition and rate will change over time, but the snapshot reveals the gap between boardroom urgency and operational deployment.

This mix of excitement and anxiety explains a lot of the noise around AI in the news, at work, and online. It is not one unified reaction. It is curiosity, genuine task-level value, competitive signalling, unresolved risk, and fear about jobs happening in the same rooms, often among the same people. Later chapters—particularly Chapter 7, *AI at Work*, and Chapter 10, *Will AI Replace Jobs?*—return to this tension directly.

Author Reflection

Before 30 November 2022, everything I knew about AI and machine learning was part and parcel of advanced tooling: classification, automation, prediction systems, monitoring, computer vision and recommendation engines.

I had used machine learning myself in systems designed to help manage infrastructure uptime and respond to problems before they became bigger ones.

It was useful technology. Interesting technology. Sometimes very clever technology.

But nobody was standing around the office proclaiming the arrival of a new machine god.

Then ChatGPT happened.

It felt a little like the opening of a *Terminator* movie, where somebody puts a date on the screen and everything after that point is suddenly different.

Almost overnight, people I knew were talking about AI. Twitter was talking about AI. The media was talking about AI. Businesses were talking about AI.

People were excited by it, frightened by it, investing in it, predicting the end of work because of it, or promising that it was about to revolutionise practically everything.

My reaction was mostly confusion.

It was like I'd spent years talking about aliens potentially visiting Earth and then one morning everyone woke up and announced:

"Have you heard about aliens?"

So I took a step back.

The hype would die down, I thought. Something else would become fashionable. Somebody would invent a new flavour of ice cream and we'd all move on.

Except that didn't happen.

AI remained.

First it was a cool technical tool. Then it became an alternative to searching Google for certain questions. Then people were using it for recipes, explanations and writing emails. Then it was helping them solve problems. Then it was their study partner. Then, occasionally, their friend.

And somewhere along the way it became:

“It’s going to take our jobs!”

Won’t somebody please think of the children?

One conversation finally made me wonder whether I was the one missing something.

A colleague asked me, “What do you think of ChatGPT?”

I shrugged.

“I honestly didn’t get why everyone was so gaga over it. It’s not like AI was new. We’d been using AI and machine learning for years, and I’d already used ML in systems to help manage uptime and respond to infrastructure problems.

“To me, ChatGPT just felt like Google and Wikipedia had been mashed together and given a chat box.

“Useful? Absolutely.

“Interesting? Definitely.

“But revolutionary?

“I didn’t see it.

“So what the hell was everyone else seeing that I wasn’t?”

He looked at me as though an alien head had just started growing out of my shoulder.

That was when I started wondering whether everybody else had drunk the Kool-Aid and I’d somehow missed getting my cup.

Then businesses everywhere seemed to discover AI at once. Construction, healthcare, education, professional services — you could barely find an industry that wasn’t discussing how AI was going to transform it.

I’m still waiting to see exactly what that transformation ultimately looks like.

Perhaps the strangest part was that I had already built conversational systems before ChatGPT.

One consumer-facing system I worked on was a Facebook chatbot designed to answer common customer questions within a defined area and hand the conversation over to a person when necessary.

Nobody thought it was miraculous.

But give people a blank text box where they can ask *their own* question, about almost anything, and suddenly the experience feels completely different.

That was probably the first clue that I had underestimated what had actually changed.

Then I Changed My Mind

Not overnight.

There wasn't some dramatic moment where I put on my nails, hair and lips and said "Here I am!". I didn't declare ChatGPT or any other AI system being hyped up as the second coming and race to join the congregation.

What changed my mind was using it.

Not playing with it for five minutes.

Actually using it.

I started giving it real problems. Things I was working on. Code. Research. Ideas. Technical questions. Writing. Problems where I already knew enough about the subject to tell when it was being useful and when it was talking absolute rubbish.

And slowly, my view of it changed.

It wasn't replacing Google. It wasn't Wikipedia with a personality. And it definitely wasn't some all-knowing machine intelligence.

What made it interesting was that I could work *with* it.

I could start with a half-formed idea and push it around until it became something useful. I could ask a follow-up question without rebuilding the entire search from scratch. I could challenge an answer, change direction, give it more context, ask it to compare approaches, find weaknesses, explain something differently or help me think through a problem.

That was different.

The conversation itself turned out to matter far more than I had given it credit for.

Traditional software usually expects you to understand how *it* wants to be used. Which screen do I need? Which menu? Which field? Which command? Which search terms will convince Google to give me what I'm actually looking for?

With conversational AI, I could start much closer to the actual problem.

I could just explain what I was trying to do.

That didn't make the machine intelligent, and it certainly didn't make it trustworthy. I found out very quickly that it could be spectacularly confident while being completely wrong.

But the more I used it, the more I realised that I'd been concentrating on the wrong question.

I had been asking:

"What's technically new about this?"

Everyone else was experiencing:

“Wait — I can just talk to it?”

Those are very different questions.

And eventually I understood why the second one mattered so much.

The underlying ideas had been developing for decades. Machine learning wasn't new. Natural-language processing wasn't new. Chatbots weren't new. Generative systems weren't entirely new either.

What changed was that all of those years of research and engineering had finally been wrapped in something almost anybody could understand immediately:

A box.

A cursor.

And the invitation to ask whatever you wanted.

The world had not suddenly discovered artificial intelligence.

For the first time, millions of ordinary people could sit down, talk directly to something that appeared intelligent, and decide for themselves what it meant.

Apparently, that makes quite a difference.

It certainly changed my mind.

None of that excitement or anxiety requires a machine mind. Language models, spam filters, recommendation systems, and fraud detectors all learn patterns, but they are not one interchangeable technology. They can use different designs, data, objectives, and checks. The family resemblance is useful for a beginner: each converts learned patterns and current input into an output. The differences matter when deciding what each system can be trusted to do.

Facts vs Reality: Sorting the Signal From the Noise

With this much money, media coverage, and workplace pressure surrounding one topic, readers need something more useful than a choice between enthusiasm and cynicism. They need a way to inspect claims.

Here are six questions worth asking whenever somebody says AI has changed everything—or that it is all nonsense:

1. **What exactly was measured?** “AI improved productivity” may mean people completed one carefully chosen writing task faster. That is evidence about the task, not every job or the entire economy.
2. **Compared with what?** A polished AI demonstration may be better than starting from a blank page but worse than an experienced professional using established tools.
3. **What does adoption mean?** Trying a free chatbot once, using it every week, paying for it, and deploying it safely across a business are four different claims.
4. **What kind of money is being counted?** Investment, company valuation, contracted spending, revenue, and profit answer different questions. None can silently stand in for the others.
5. **Who supplied the number?** A company can be the best source for the date and terms of its own product launch. It is not an independent judge of how revolutionary or safe that product is.
6. **Is this an observation or a forecast?** “This system completed this task” describes evidence. “This will eliminate half of all jobs” describes a model of the future, with assumptions that deserve inspection.

These questions do not make every claim false. They make the claim small enough to evaluate.

What We Can Say Honestly

We can say that modern generative AI has real and unusual capabilities. It can produce and transform language, images, audio, video, and code through a general-purpose interface. Research has measured meaningful benefits in particular tasks. Usage has grown quickly. Investment is enormous. Businesses and governments are taking the technology seriously.

We can also say that fluent output remains unreliable, benefits vary by task and user, adoption is uneven, and long-term effects on organisations, jobs, education, culture, and productivity remain uncertain. Major institutions studying the subject say so plainly. That is not anti-AI pessimism. It is the current evidence.

History cuts both ways. Earlier AI booms show that technical progress can be real while commercial promises outrun it. They also show that disappointment with one approach does not end the field. Expert systems lost their starring role, but AI research continued. A future correction in investment or valuations would not prove that generative AI was useless. Continued technical progress would not prove that every current company, product, or forecast was sound.

This Book's Position

This series has an opinion, and it is not “both sides have a point” for the sake of sounding safe.

The useful position is that AI is already capable enough to learn, use, and scrutinise. It is also unreliable enough that judgement, verification, and responsibility cannot be outsourced to it. The commercial race is powerful enough to shape the tools placed in front of us, but commercial confidence is not scientific proof. Anyone demanding that you either worship the technology or dismiss it entirely is asking you to stop thinking at the moment thinking becomes most valuable.

The aim of this series is therefore practical AI literacy. You should be able to use a tool without accepting its marketing story, recognise a genuine result without turning it into a universal promise, and reject nonsense without pretending the underlying technology has no value. That is how we will separate the truth from the BS throughout the chapters ahead.

Core Takeaway

AI is not new, and there was no single moment when it woke up. What changed was a convergence: stronger generative models, a conversational interface, cheap public access, vast investment, product distribution, and intense publicity. Conversation made the technology feel personal; generation made it broadly useful; money and competition made it difficult to avoid.

Recap

The AI in Chapter 1 often predicts narrow answers quietly in the background. The AI that became publicly visible after 2022 can generate open-ended content in conversation. It did not arrive from nowhere, and access was not the only thing that improved. Years of research, computing infrastructure, investment, and product design met a simple interface at the same time. Treat demonstrated capability as real, forecasts as forecasts, and financial enthusiasm as evidence of a bet—not proof that the bet will win.

What AI Can Actually Do

By now, you know that AI finds patterns, makes predictions, and can generate new material. That is useful background, but it still leaves the question most people actually care about:

What can I do with it?

The honest answer is both more impressive and less magical than the advertising suggests. AI can help with a surprisingly wide range of tasks. It can turn rough notes into an email, shorten a long document, explain an unfamiliar subject, generate a picture, analyse comments, translate a message, help write software, and carry out parts of a repeated workflow.

It can also do one task brilliantly and then stumble over another that looks easier. Its abilities are broad, but they are not even. Learning to use AI is therefore not about discovering one giant button labelled **Do Everything**. It is about recognising the kinds of work where AI gives you a useful advantage, and knowing what your part of the job becomes afterwards.

Key Idea

AI is most useful when it changes the cost of getting from an idea, question, or pile of information to a workable first version.

That first version might be a paragraph, a plan, an image, an explanation, a set of options, or the beginning of an automated process. It does not need to be perfect to be useful. If it saves you from starting with a blank page—or helps you see something you would otherwise have missed—it has already changed the task.

Four Useful Ways to Think About AI

A long list of AI features quickly becomes overwhelming. It also dates badly, because individual products keep adding, renaming, and removing features. A more durable way to understand AI capability is through four ordinary verbs: **generate, transform, interpret, and act**.

To **generate** is to make a new first version. You might ask for an email, ten gift ideas, a picture of a fictional place, a set of quiz questions, or a small piece of computer code.

To **transform** is to take something that already exists and change its form. Rough notes become polished prose. A long report becomes a short summary. A formal announcement becomes a friendly message. Spoken words become a transcript. English becomes Spanish, Japanese, or Arabic.

To **interpret** is to examine material and suggest what it means. AI can group customer comments into themes, compare two documents, explain a chart, describe an image, identify recurring topics, or help you understand code someone else wrote.

To **act** is to use tools and perform steps rather than merely produce an answer. An AI system might search several sources, update a file, prepare a weekly report, or move information between applications. This is where people often use words such as *automation* or *agent*. It is also the category where a useful assistant can create the most trouble if it is given too much freedom without checks.

These categories overlap. Asking AI to turn a voice note into a task list uses voice recognition, transformation, and interpretation. Asking it to research holiday options and place them in a comparison table combines generation, interpretation, and possibly action. The categories are not boxes that every task must fit into perfectly. They are handles you can use when thinking, “Could AI help with this?”

A particularly good starting question is:

Is this work expensive for me to produce, but reasonably easy for me to check?

If the answer is yes, you may have found a strong AI task.

Writing: Escaping the Blank Page

Writing is one of the most immediately useful applications because so much everyday and professional life begins with words. Emails, letters, complaints, applications, meeting notes, reports, social posts, invitations, product descriptions, and awkward messages all require someone to create a first sentence.

AI is very good at giving you something to react to.

Suppose you have these notes:

- project is late
- customer is waiting
- need a new delivery date
- ask for a reply today
- keep it polite

An AI tool can turn them into a structured email in seconds. You can then say, “Make it warmer,” “Shorten it,” or “Be firmer about needing an answer today.” You are not asking the tool to discover the facts. You are supplying the facts and asking it to shape them.

This is more than a party trick. In a controlled 2023 experiment involving mid-level professional writing tasks, people with AI access completed the work faster and produced outputs that were rated more highly on average. The study does not prove that AI improves every kind of writing, but it does show that the blank-page benefit can produce measurable gains in the right setting.¹

The most useful shift is often not from “human writing” to “machine writing.” It is from **human produces everything** to **AI drafts, human directs and edits**. Instead of spending ten minutes trying to find an opening sentence, you spend those minutes deciding which opening works and making it sound like you.

That distinction matters. A polished first draft can still contain the wrong date, the wrong tone, or a promise you never intended to make. If you ask AI to write a complaint to an airline, it may produce confident language about consumer law that sounds impressive but does not apply where you live. The writing may be excellent while the advice inside it is poor.

Use AI to reduce the mechanical effort of writing, not to surrender ownership of what the words say.

Summarising: Turning More Into Less

Summarising is another natural fit. Modern life produces far more text than most people have time to read: long email chains, meeting transcripts, policies, research papers, school notices, reports, terms and conditions, and group chats that somehow accumulated 147 messages while you were making dinner.

AI can condense that material into a shorter form. You can ask for five key points, a list of decisions, unanswered questions, action items, or an explanation written for someone unfamiliar with the subject.

The safest version of this task has three features:

1. You provide the source material.
2. You tell the AI what kind of summary you need.
3. You can compare the result with the original.

That is different from asking, “What did this report say?” without supplying the report. In the first case, the evidence is in front of the model. In the second, you are asking it to recall or infer material that may not be available to it at all.

Even source-based summaries need checking. AI can give too much weight to a vivid detail, miss an exception buried near the end, or write a neat conclusion the original author never made. The shorter the summary, the more judgement is involved in deciding what to leave out.

This does not make summarising useless. It makes the output a map rather than the territory. A summary helps you decide where to focus. If the document matters, you still return to the source before relying on a claim.

Try This

Choose an article, email, or document you already understand. Ask an AI tool to summarise it in five bullet points. Check every point against the original. Then ask for the same material to be explained to a teenager or turned into three action items. Notice both what the AI preserves and what it changes.

Brainstorming: More Doors to Open

Brainstorming is not the same as asking AI for the correct answer. It is asking for possibilities.

You might request ten low-cost birthday ideas, names for a gardening club, different ways to organise a presentation, meal ideas based on what is in the fridge, or questions to ask before accepting a job. The first suggestion may be obvious. The third may be strange. The seventh may prompt an idea neither you nor the AI stated directly.

That is the value: AI gives your mind more material to push against.

Research on AI-assisted creativity illustrates an interesting trade-off. Some experiments have found that people using AI produce ideas rated as more creative. Other work has found that when many people rely on similar AI suggestions, the collection of ideas can become less diverse. Individual results may improve while the room as a whole starts thinking in the same direction.²

The practical response is not to avoid AI. It is to ask it to widen the search rather than close it too early. Instead of “Give me the best idea,” try:

- “Give me ten genuinely different directions, not ten versions of one idea.”
- “What assumptions are all of these options making?”
- “Suggest three approaches that avoid the obvious solution.”
- “Now argue against your own suggestions.”
- “How would three very different kinds of people approach this?”

AI is useful here as an option generator, critic, and provocateur. It is less useful as the unquestioned judge of which option deserves to win.

Images: Making an Idea Visible

Image generation turns a written description into a visual starting point. You can describe a fictional book cover, a room layout, a character, a birthday invitation, a product concept, or a scene for a presentation and receive several interpretations within moments.

For a beginner, the biggest benefit is often not producing a finished work of art. It is making an idea visible enough to discuss.

Imagine trying to explain a home-office redesign. “Warm, simple, lots of natural timber, green walls, room for two monitors, but not like a corporate office” is understandable, but still abstract. A set of generated concepts lets you point and say, “This colour, but not that desk,” or “I like the lighting, but the room feels too expensive.” You have moved from imagining alone to reacting to something concrete.

The same applies to storyboards, event themes, packaging ideas, garden layouts, and fictional scenes. Generated images can help explore direction before paying for production or asking a designer to create a finished piece.

That does not erase the value of photographers, illustrators, designers, or artists. A professional brings judgement, taste, technical craft, and an understanding of the purpose behind the image. AI can cheaply generate possibilities. Deciding what communicates well—and whether an image is appropriate to publish—is a different job.

Copyright, consent, deepfakes, and style imitation deserve much more attention than this chapter can give them. For now, the useful distinction is simple: generating a visual concept is not automatically the same as creating a finished asset you are entitled to use anywhere.

Code: Useful, Testable, and Uneven

You do not need to be a programmer to understand why coding is such a useful example of AI capability.

Code is language written under strict rules. It also gives us unusually clear ways to test whether an answer works. A programmer can ask AI to explain an error, draft a small function, generate test cases, translate code from one programming language to another, or identify unusual cases the original author forgot to consider.

Sometimes this produces dramatic savings. Sometimes reviewing and correcting the AI takes longer than writing the code directly. Studies of AI-assisted software development have reported both gains and slowdowns in different settings. In one realistic study, experienced open-source developers working in large projects they already knew took longer when early-2025 AI tools were available, even though they believed the tools had made them faster. The researchers were careful not to claim that AI slows all developers; their finding applied to a particular group, tools, and kind of work.³

That is the broader lesson. “Can AI write code?” is too simple a question. It certainly can. Better questions are:

- Does the person using it understand the goal?
- Can the result be tested?
- How expensive is a hidden mistake?
- Does the AI have enough context about the existing system?
- Is reviewing the generated code easier than writing it?

AI can be an excellent coding assistant. Generated code is still code: it needs to be understood, tested, and maintained after the exciting demonstration is over.

Translation: Fluent Is Not Always Exact

AI translation can be extraordinarily convenient. It can help you understand a menu, exchange messages with a relative, read the basic meaning of a notice, or produce a first draft in another language. It can also follow instructions about tone: formal or informal, literal or natural, suitable for a customer or suitable for a friend.

That is a meaningful improvement over replacing one word at a time. Languages do not line up like columns in a spreadsheet. Tone, context, idiom, humour, and cultural expectations all affect what a sentence means.

Quality still varies by language, subject, and the material available to the model. A translation can sound beautifully fluent while slightly changing the meaning. Specialist terminology can be mishandled. Less widely represented languages may receive weaker results than languages with enormous amounts of available text.⁴

For everyday understanding, AI translation may be more than enough. For a contract, medical instruction, public statement, or anything else where exact wording carries consequences, fluent is not the same as professionally verified.

One helpful technique is to ask for both a literal translation and a natural conversational version, followed by an explanation of any phrase that could be interpreted in more than one way. This does not guarantee correctness, but it exposes some of the choices that a single smooth translation can hide.

Tutoring: A Conversation That Can Change Direction

Search is excellent when you know what to look for. Learning is harder when you do not yet understand enough to ask the right question.

AI can explain a concept, respond to a follow-up question, change the level of difficulty, invent another example, quiz you, or ask you to explain the idea back in your own words. If the first explanation does not land, you can say, “Try again without the technical terms,” or “Show me with an example involving football.” The conversation can change direction each time you respond.

That adaptability is valuable. A carefully designed AI tutor used in a 2025 randomised study of an undergraduate physics course produced strong learning results in that particular setting. The important detail is that it was not simply a generic chatbot told to teach physics. It was built around the course content and established teaching practices, with structured activities and prepared solutions.⁵

The sensible conclusion is not that AI is better than teachers. It is that well-designed AI tutoring can be effective, and that ordinary users can borrow some of the same principles.

Instead of asking, “Teach me economics,” try:

“Explain inflation to me as a beginner. Use one everyday example, then ask me one question at a time to check my understanding. If I get something wrong, give me a hint before giving me the answer.”

That request asks for a learning process, not just an information dump.

AI can be patient, available, and willing to rephrase something twelve times. It does not replace the experience and judgement of a good teacher, and it can teach incorrect information if left unchecked. But it can provide a low-pressure space to ask the question you were embarrassed to ask in front of a room.

Analysis: Looking for Shape in the Material

Analysis is where AI begins to feel less like a writing tool and more like a thinking companion.

You can give it customer comments and ask for recurring themes. You can provide a table of household spending and ask which categories changed most. You can compare two versions of a policy and ask what was added or removed. You can supply survey responses and request a first attempt at grouping them.

The best analysis prompts keep the AI close to evidence you can inspect. For example:

“Group these twenty comments into themes. Under each theme, quote the comments that support it. Put uncertain comments in a separate group rather than forcing them to fit.”

That is better than “Tell me what customers think,” because it asks the tool to show its working material. You can see whether the theme actually matches the comments.

Analysis does not mean handing a spreadsheet to AI and treating every conclusion as fact. The tool may mistake correlation for cause, overlook missing data, or invent a tidy explanation for an unusual number. It is often good at finding a possible pattern. Establishing that the pattern is real—and deciding what it means—still requires judgement.

The aim is not to remove thought. It is to move some effort away from sorting and towards questioning.

Voice: Another Way In and Out

Voice is not one separate intelligence. It is another way to interact with AI.

Speech recognition can turn spoken words into text. A language model can work with that text. A speech system can then read the response aloud. Combined, these abilities make AI feel like a conversation rather than a form.

That can be useful when your hands are busy, typing is difficult, or speaking is simply the fastest way to capture an idea. A five-minute rambling voice note can become a transcript, a short summary, and a list of next steps. An unfamiliar name or heavy background noise may be transcribed incorrectly, but correcting two words can still be faster than typing the whole note.

Voice can also make AI more accessible, but it changes the privacy situation. Speaking sensitive information aloud in a public place or sending a private recording to an online service deserves the same thought as pasting confidential text into a chat box. Convenience does not make the information less personal.

Automation: From Answering to Doing

The capabilities so far mostly produce something for you to review. Automation goes a step further by connecting the AI to tools that can take action.

A simple workflow might collect a set of weekly updates, group them by project, and draft a report. Another might turn meeting notes into proposed tasks. A more advanced system might search websites, edit files, run software, fill in forms, or coordinate several steps towards a goal.

This can save enormous amounts of repetitive effort. It also changes the risk. A poor draft is annoying. A poorly checked action can send the wrong email, modify the wrong file, spend money, or publish something that was never ready.

A useful way to introduce automation is to begin with **prepare, do not commit**. Let the AI draft the report but not send it. Let it propose file changes but require approval before saving them. Let it prepare a booking comparison but not purchase the ticket. As confidence grows and the task becomes well understood, some low-risk steps can be allowed to run automatically.

Autonomy and reliability are not the same thing. A system being capable of taking ten steps does not mean it will take all ten correctly every time.

The Jagged Frontier

People often talk about AI as though it has one level of intelligence. They ask whether it is smart, whether it is smarter than a person, or whether the newest system is smarter than the previous one.

Real capability does not form a smooth ladder. It looks more like a jagged skyline: very high in some places, unexpectedly low in others.

Researchers studying professionals using AI described this as a **jagged technological frontier**. On tasks that sat inside the system's abilities, AI users worked faster and produced stronger results. On a task outside that frontier, AI could steer people towards the wrong conclusion.⁶

This explains an experience that confuses many new users. One minute, an AI tool rewrites a difficult letter beautifully. The next, it miscounts items in a simple list. It can explain an advanced concept in clear language and then invent a book title. It can produce working code for an unfamiliar task and fail to follow an apparently simple instruction in a document.

Difficulty for AI is not identical to difficulty for humans. We cannot assume that because it handled Task A, it will handle Task B. We also cannot assume that because it failed once, it is useless for everything else.

The frontier moves as tools improve, but it remains jagged. That is why learning to use AI is partly a process of experimentation. Try it on a low-risk task. Inspect the result. Learn where it helps. Add checks where it fails. Do not grant broad trust merely because one result was impressive.

How a Language Model Builds an Answer

To understand both the usefulness and the mistakes, it helps to look underneath the conversation for a moment.

A language model does not begin with a shelf of complete answers and choose one. It builds a response piece by piece. Text is broken into small units called **tokens**—sometimes a word, sometimes part of a word, and sometimes punctuation. At each step, the model uses patterns learned during training plus the context of your current request to estimate plausible next tokens. The selected token becomes part of the context for the following step, and the response continues.⁷

The model does not always choose the single most probable continuation. Depending on how the system is configured, it may select among several likely possibilities. That variation is one reason the same prompt can produce different wording on different attempts.

For a beginner, the most important point is this:

The answer is generated, not retrieved as a guaranteed finished fact.

That process is remarkably flexible. It is why the same kind of system can draft a poem, explain a tax term, restructure notes, translate a message, or role-play an interviewer. It has learned a vast number of relationships in language and can continue a pattern in ways that respond to your instructions.

But fluent continuation is not the same as fact checking. The system can generate the shape of a correct answer—the confident tone, tidy explanation, plausible names, and formal-looking citations—even when a crucial detail is wrong.

Why Confident Wording Can Still Be Wrong

Humans use tone as evidence. A person who hesitates may be unsure. A person who answers smoothly and precisely sounds as if they know what they are talking about.

Language models break that familiar connection. Their job is to produce language that fits. Smooth delivery is therefore a feature of the output, not proof that the underlying claim has been checked.

This is why an invented date can arrive in exactly the same tone as a real one. It is why a fabricated citation can look properly formatted. It is why asking the same question again can produce a different number without any visible change in confidence.

Researchers and standards bodies use terms such as *hallucination* and *confabulation* for plausible but false generated content. Work on this problem also shows that systems can be pushed towards guessing when evaluations reward a correct answer but do not reward an honest “I don’t know.” Better training, retrieval, tools, and verification can reduce the problem, but no polished tone should be mistaken for a built-in truth detector.⁸

Watch Out

Tone is not evidence. The myth is that a confident answer is probably a checked answer. The reality is that AI can produce confident wording whether a detail is correct, mistaken, or invented. Check the claim, not the confidence.

This does not mean you must treat every AI response as worthless until proven in court. Verification should match the stakes. A silly name for a family group chat needs less checking than medical advice. A draft invitation needs less checking than a contract. A list of dinner ideas needs less checking than instructions involving money, safety, or another person's rights.

Usefulness does not require perfection. It requires knowing what kind of error you can tolerate.

Your Job Changes; It Does Not Disappear

AI is often described as removing work. A better description is that it changes the division of work.

The tool can generate, transform, and sort quickly. You provide the goal, the relevant context, the boundaries, and the decision about what is good enough. You keep responsibility for facts that matter and actions that have consequences.

That gives you a simple working loop:

1. **Direct** — explain the goal and supply useful context.
2. **Inspect** — read or test what the AI produced.
3. **Select** — keep the useful parts and reject the rest.
4. **Verify** — check important claims against reliable evidence.
5. **Refine** — edit the output or ask for another attempt.

Research with knowledge workers suggests that AI can shift critical thinking towards verification, integration, and stewardship. It also found an association between greater confidence in AI and less reported critical thinking, which is a useful warning against letting convenience turn into automatic trust.⁹

The best result is not always the workflow with the least human effort. It is the workflow that spends human effort where it matters most.

The author reflection can then ground that workflow in lived experience:




Figure 3-1: Useful AI work is a loop: your goal and context go to the AI, which generates, transforms, interprets or acts; you judge the output by keeping, editing, verifying or rejecting it, then either use the result or refine and send it back to the AI.

Author Reflection

One of the better examples of AI genuinely saving me time came from rebuilding a piece of software I had already written.

The application is called **CurseDelete**. Its job is fairly simple to describe: delete very large directory trees quickly, including the sort containing hundreds of thousands or millions of files where the normal operating-system tools can become painfully slow, run into permissions problems or generally decide they have had enough for the afternoon.

The original version was written in C#.

It already worked across Windows, Linux and macOS, and it was reasonably clever about how it did the job. Instead of discovering every file first and then starting the deletion, it used a bounded producer-and-consumer pipeline so enumeration and deletion happened at the same time. That kept memory usage controlled and allowed multiple workers to delete files in parallel.

It worked.

But eventually I wanted to build it properly for what it had become.

I wanted something closer to the operating system rather than sitting several layers above it. I wanted better control over filesystem operations, better performance characteristics, stronger safety guarantees and, importantly, I wanted to stop pretending that Windows, Linux and macOS were basically the same operating system with different path separators.

They are not.

That way lies sadness.

So I started rebuilding it from scratch in Rust as **CurseDelete 2**.¹⁰

This is where AI saved me a considerable amount of time.

I already understood the application, the concurrency model and what I wanted the software to do. What I did not have sitting conveniently in my head was every detail of Darwin filesystem APIs, Linux filesystem behaviour, Windows native deletion semantics and the various delightful little differences operating systems have accumulated over several decades specifically to ruin a programmer's afternoon.

AI became an extremely useful research and engineering partner.

Instead of spending hours moving between API documentation, old Unix references, Stack Overflow posts and kernel documentation just to establish which primitives I should investigate, I could ask very specific questions about the problem I was solving.

For example, on macOS and Linux I moved towards native POSIX operations such as `openat`, `unlinkat` and `fstatat`, rather than repeatedly resolving complete path strings through higher-level filesystem APIs.

That was not merely about making it faster.

It also exposed a security issue I wanted the new version to handle better. If you discover a filesystem object using one path and later resolve that entire path again when you delete it, something can potentially change underneath you between those two operations. Filesystem people have given this the wonderfully friendly name **time-of-check to time-of-use**, or TOCTOU, because apparently even race conditions deserve an acronym.

Using directory file descriptors and relative operations can reduce that window considerably.

Linux then introduced its own collection of details.

Creation timestamps are a good example. macOS exposes a birth time in its normal filesystem metadata. Traditional Linux `stat` does not. Linux instead has `statx`, where supported, which can expose a birth timestamp depending on the kernel and filesystem.

That matters if your deletion tool supports something like:

“Delete files more than 30 days old.”

Suddenly the definition of *old* is not quite as universal as the command line makes it look.

AI was extremely useful for getting me rapidly to those questions.

But it could not make the engineering decisions for me.

I still had to decide what `CurseDelete` should do when a timestamp does not exist. In this case, the safe answer was not to shrug, guess that the file was probably ancient and enthusiastically delete it.

If the required timestamp cannot be established, retain the file.

Boring.

Safe.

Exactly what I want from software whose primary feature is destroying things.

The same thing happened with performance.

Linux exposes lower-level facilities such as `getdents64`, and it would have been very easy to convince myself that dropping down another layer must automatically make `CurseDelete` faster.

It sounds wonderfully hardcore.

The problem is that `readdir` is already backed by those mechanisms, and I did not have evidence that directory-entry parsing was actually the bottleneck.

So I didn’t implement the cleverer-looking answer.

The rule became: **benchmark first**.

If the numbers eventually show that the extra complexity buys something meaningful, then I can change it.

That distinction became important throughout the project.

AI could very quickly tell me what was possible.

My job was deciding what was justified.

It also helped me understand where the operating systems genuinely deserved different implementations rather than forcing everything through one lowest-common-denominator abstraction.

On Linux, for example, CurseDelete 2 can inspect `/proc/[pid]/fd` to discover which processes hold files open. That means it can do the job directly through facilities Linux already exposes.

macOS does not provide the same interface, so the implementation uses `lsof`.

Windows is different again, with facilities such as Restart Manager, Windows security APIs and its own native file-deletion mechanisms.

That became one of the architectural decisions I am happiest with.

The shared part of CurseDelete understands things such as streaming deletion, filtering, retention rules, worker control, reporting and safety policy.

The platform engines understand the operating systems.

That seems obvious when written down.

It took considerably more work to get there.

AI shortened that work enormously because I could use it to explore APIs, compare approaches, challenge designs, generate implementation candidates and build test cases without beginning every unfamiliar operating-system detail from a blank search box.

But I still had to read the code.

I still had to understand why it worked.

I still had to decide whether a proposed abstraction actually belonged in the architecture.

And because this particular application deletes files for a living, “it compiled” was never going to qualify as sufficient testing.

One particularly useful lesson was that AI often got me to **80 or 90 per cent of the answer extraordinarily quickly**.

That final ten per cent was where most of the engineering lived.

Does this behave the same way on the other operating system?

What happens when a symlink appears?

Can this accidentally resolve outside the requested directory?

What happens when a timestamp is unavailable?

Does increasing the worker count actually make deletion faster, or have I merely invented a very efficient way of making the storage device miserable?

And, perhaps most importantly:

What happens when this thing is wrong?

CurseDelete 2 ended up being much more than a Rust translation of the C# version.

The new version has native platform engines, adaptive concurrency, structural protection against deleting filesystem roots, safer symlink handling, retention and filtering, machine-readable output and a clearer separation between the common deletion engine and the operating-system-specific machinery underneath it.

AI helped me get there much faster.

But the interesting part is *why* it saved time.

It did not replace the bit where I understood my own software.

It removed a large amount of the mechanical work between “**I need to understand how Linux does this**” and “**here are the relevant mechanisms, trade-offs and several possible implementations.**”

That left me spending more of my time on the part I actually wanted to be doing: deciding which implementation was correct for the product, testing the dangerous bits and occasionally telling the AI that, while its solution was technically impressive, we did not need to rebuild half the operating system to delete a file.

That, for me, is where AI-assisted software development becomes genuinely useful.

Not when I can generate code faster.

When I can get from a problem I understand to a solution I trust in substantially less time.

Core Takeaway

AI can do a great deal: write, summarise, brainstorm, generate images, help with code, translate, tutor, analyse, work through voice, and automate parts of a process. Those abilities become easier to remember when you think in four verbs: generate, transform, interpret, and act.

Its capability is also jagged. An impressive result in one area does not prove reliability in the next. A language model builds its answer piece by piece from learned patterns; it does not retrieve a guaranteed complete answer. That is why flexibility and error arrive together.

Your role is not merely to press a button. It is to direct, inspect, select, verify, and refine. AI can perform the fast first pass. You decide what deserves to survive it.

Recap

AI is a broad but uneven assistant. It is particularly useful for generating, transforming, and interpreting information, and it can increasingly act through other tools. Start with tasks that are expensive to produce but easy to check. Remember that fluent language is generated, not guaranteed, and keep human judgement where mistakes matter.

Notes

Chapter 02: Why Everyone Suddenly Talks About AI

- 1 OpenAI, “Introducing ChatGPT,” 30 November 2022, <https://openai.com/index/chatgpt/>. The launch post described ChatGPT as a free research preview and listed plausible-sounding incorrect answers among its limitations.
- 2 Daniel Jurafsky and James H. Martin, *Speech and Language Processing*, 3rd ed. online manuscript (2026), <https://web.stanford.edu/~jurafsky/slp3/>.
- 3 John McCarthy, Marvin L. Minsky, Nathaniel Rochester, and Claude E. Shannon, “A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence,” dated 31 August 1955, *AI Magazine* 27, no. 4 (2006), <https://doi.org/10.1609/aimag.v27i4.1904>.
- 4 UK Parliament, House of Lords Artificial Intelligence Committee, “AI in the UK: Ready, Willing and Able?” 2018, section “AI through the ages,” <https://publications.parliament.uk/pa/ld201719/ldselect/ldai/100/10018.htm>.
- 5 Defense Advanced Research Projects Agency, “Artificial Intelligence,” in *DARPA: 60 Years 1958–2018*, pp. 96–103, <https://www.darpa.mil/sites/default/files/attachment/2025-02/magazine-darpa-60th-anniversary.pdf>.
- 6 Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton, “ImageNet Classification with Deep Convolutional Neural Networks,” *Advances in Neural Information Processing Systems* 25 (2012), <https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-networks>; Ashish Vaswani et al., “Attention Is All You Need,” arXiv:1706.03762 (2017), <https://arxiv.org/abs/1706.03762>.
- 7 OpenAI, “Introducing ChatGPT,” 30 November 2022, <https://openai.com/index/chatgpt/>. The launch post described ChatGPT as a free research preview and listed plausible-sounding incorrect answers among its limitations.

- 8 Chloe Autio et al., *Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile*, NIST AI 600-1 (2024), <https://doi.org/10.6028/NIST.AI.600-1>.
- 9 Stanford Institute for Human-Centered Artificial Intelligence, “Economy,” *2025 AI Index Report*, <https://hai.stanford.edu/ai-index/2025-ai-index-report/economy>. Its categories and methodology should be checked again before a later print edition updates the figures.
- 10 Microsoft, “OpenAI Forms Exclusive Computing Partnership with Microsoft to Build New Azure AI Supercomputing Technologies,” 22 July 2019, <https://news.microsoft.com/source/2019/07/22/openai-forms-exclusive-computing-partnership-with-microsoft-to-build-new-azure-ai-supercomputing-technologies/>; Microsoft, “Microsoft and OpenAI Extend Partnership,” 23 January 2023, <https://blogs.microsoft.com/blog/2023/01/23/microsoftandopenaiextendpartnership/>; Amazon, “Amazon and Anthropic Deepen Their Shared Commitment to Advancing Generative AI,” 27 March 2024, <https://www.aboutamazon.com/news/company-news/amazon-anthropic-ai-investment>.
- 11 Olivia Sidoti and Colleen McClain, “34% of U.S. Adults Have Used ChatGPT, About Double the Share in 2023,” Pew Research Center, 25 June 2025, <https://www.pewresearch.org/short-reads/2025/06/25/34-of-us-adults-have-used-chatgpt-about-double-the-share-in-2023/>.
- 12 Flavio Calvino, Jelmer Reijerink, and Lea Samek, “The Effects of Generative AI on Productivity, Innovation and Entrepreneurship,” OECD Artificial Intelligence Papers, no. 39 (2025), <https://doi.org/10.1787/b21df222-en>.
- 13 Kathryn Bonney et al., “Tracking Firm Use of AI in Real Time: A Snapshot from the Business Trends and Outlook Survey,” US Census Bureau Center for Economic Studies Working Paper CES-24-16 (2024), <https://www.census.gov/library/working-papers/2024/adrm/CES-WP-24-16.html>.

This chapter was developed from the author’s viewpoint with research and drafting assistance from ChatGPT and Codex. The sources below support its research-dependent claims. Company announcements are used for the terms of their own launches and partnerships, not as independent proof of those products’ value.

Chapter 03: What AI Can Actually Do

- 1 Shakked Noy and Whitney Zhang, “Experimental Evidence on the Productivity Effects of Generative Artificial Intelligence,” *Science* 381, no. 6654 (2023), 187–192, <https://doi.org/10.1126/science.adh2586>.
- 2 Byung Cheol Lee and Jaeyeon Chung, “An Empirical Investigation of the Impact of ChatGPT on Creativity,” *Nature Human Behaviour* (2024), <https://doi.org/10.1038/s41562-024-01953-1>; Lennart Meincke, Gideon Nave, and Christian Terwiesch, “ChatGPT Decreases Idea Diversity in Brainstorming,” *Nature Human Behaviour* 9 (2025), 1107–1109, <https://doi.org/10.1038/s41562-025-02173-x>.
- 3 Joel Becker et al., “Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity,” METR (2025), <https://metr.org/blog/2025-07-10-early-2025-ai-experienced-os-dev-study/>. The researchers explicitly caution against generalising the result to all software development.
- 4 Jianhao Yan et al., “GPT-4 vs. Human Translators,” arXiv (2024), <https://arxiv.org/abs/2407.03658>; Wenhao Zhu et al., “Multilingual Machine Translation with Large Language Models,” *Findings of NAACL 2024*, <https://aclanthology.org/2024.findings-naacl.176/>.

- 5 Greg Kestin et al., “AI Tutoring Outperforms In-Class Active Learning: An RCT Introducing a Novel Research-Based Design in an Authentic Educational Setting,” *Scientific Reports* 15 (2025), article 17458, <https://doi.org/10.1038/s41598-025-97652-6>.
- 6 Fabrizio Dell’Acqua et al., “Navigating the Jagged Technological Frontier,” Harvard Business School AI Institute (2023), <https://aiinstitute.hbs.edu/navigating-the-jagged-technological-frontier/>.
- 7 Daniel Jurafsky and James H. Martin, *Speech and Language Processing*, 3rd ed. online manuscript (2026), <https://web.stanford.edu/~jurafsky/slp3/>.
- 8 Chloe Autio et al., *Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile*, NIST AI 600-1 (2024), <https://doi.org/10.6028/NIST.AI.600-1>; Adam Tauman Kalai et al., “Why Language Models Hallucinate,” OpenAI (2025), <https://openai.com/index/why-language-models-hallucinate/>.
- 9 Hao-Ping Lee et al., “The Impact of Generative AI on Critical Thinking: Self-Reported Reductions in Cognitive Effort and Confidence Effects From a Survey of Knowledge Workers,” *CHI 2025*, <https://doi.org/10.1145/3706598.3713778>.
- 10 RePass Cloud, “CurseDelete 2,” product page, <https://repasscloud.com/products/cursedelete/>. Checked 3 October 2026: describes CurseDelete 2 (version 2.0.0, released 17 August 2026) as a Rust rewrite with a command-line interface that streams deletion while it enumerates, keeps memory use bounded and protects filesystem and share roots, with separate engines for macOS, Linux and Windows. The page offers a free Community edition for personal use and paid Business and Enterprise editions. This is the author’s own product, sold by the publisher of this book. The technical details in the reflection describe the author’s design, as supplied by the author: the original C#/.NET version used a bounded producer-and-consumer pipeline; version 2 shares one core engine and adds platform engines that use directory-relative POSIX calls (`openat`, `unlinkat`, `fstatat`) on macOS and Linux, `statx` birth times on Linux where the kernel and filesystem provide them, `/proc/[pid]/fd` on Linux and `ls` on macOS to find open files, and native Windows deletion, security and Restart Manager interfaces. According to the author, the project documentation records that these measures reduce but do not remove the time-of-check to time-of-use risk, and that lower-level directory reading was not adopted without benchmarks. Product details and prices may change.

This chapter was developed from the author’s viewpoint with research and drafting assistance from ChatGPT and Codex. The sources below informed the factual claims and editorial framing; the chapter synthesises them in its own words. The author reflection is the author’s own account of rebuilding CurseDelete, the author’s own software, and was used word for word from the author’s comment on GitHub issue #21. CurseDelete 2 is sold by RePass Cloud Pty Ltd, which also publishes this book; it is named here at the author’s request and is not a recommendation.

Further Reading

Every source this book relies on is listed in the Notes, chapter by chapter. That list is long, and much of it is research papers, court records and technical reports. You don't need to read any of it to use AI well.

This page is shorter. It picks out the sources worth your time if you want to go a little further, mostly free, plain-language pages from Australian government agencies, plus a few reports the book leans on. Each entry says which chapter it connects to.

Web pages move and AI products change quickly. Everything here was checked while the book was being written, in 2026. If a link no longer works, search for the title. The companion website keeps an up-to-date version of this list.

Understanding What AI Can and Can't Do

- OpenAI Help Center, “Does ChatGPT Tell the Truth?” <https://help.openai.com/en/articles/8313428-does-chatgpt-tell-the-truth>. A short page from a chatbot's own maker explaining why it can sound sure of itself and still be wrong. It goes with Chapter 5.
- Stanford Institute for Human-Centered Artificial Intelligence, *The 2026 AI Index Report*, <https://hai.stanford.edu/ai-index/2026-ai-index-report>. A large yearly report on where AI actually stands: what it can do, who is building it, and how it is being used. Dip into the chapter summaries rather than reading it cover to cover. It goes with Chapters 9 and 11.

Using AI at Home and at Work

- ASIC Moneysmart, “AI and money decisions,” <https://moneysmart.gov.au/online-safety/ai-and-money-decisions>. Practical advice on using AI for money questions: share only what the task needs, and check what it tells you before you act. It goes with Chapter 6.
- OpenAI Help Center, “Data controls in ChatGPT,” <https://help.openai.com/en/articles/7730893-data-controls-in-chatgpt>, and Google, “Gemini Apps Privacy Hub,” <https://support.google.com/gemini/answer/13594961>. Where to find the settings that decide what happens to your conversations. Other AI tools have similar pages. They go with Chapters 9 and 13.
- Office of the Australian Information Commissioner, “Australian Privacy Principles quick reference,” <https://www.oaic.gov.au/privacy/australian-privacy-principles/australian-privacy-principles-quick-reference>. A one-page summary of the rules Australian organisations follow with your personal information. Useful before you paste anything personal into an AI tool at work. It goes with Chapter 13.

Jobs, Skills and Hype

- Jobs and Skills Australia, *Australia’s AI Transition: Jobs, Skills and the Future of Work*, <https://www.jobsandskills.gov.au/studies/generative-artificial-intelligence-capacity-study>. The Australian Government’s study of how generative AI is likely to affect work, looking at tasks rather than whole job titles. It goes with Chapters 10 and 12.
- Joel Becker, Nate Rush, Beth Barnes and David Rein, “Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity,” METR (10 July 2025), <https://metr.org/blog/2025-07-10-early-2025-ai-experienced-os-dev-study/>. A careful study in which experienced programmers felt faster with AI but, on the measured tasks, were slower. It’s a narrow, early-2025 result, not a verdict on AI in general, and it’s a good reminder to measure rather than assume. It goes with Chapters 3, 7, 10 and 11.
- Center for AI Safety, “Statement on AI Risk” (30 May 2023), <https://safe.ai/work/statement-on-ai-extinction-risk>. The one-sentence statement on extinction risk discussed in Chapter 11, with its list of signatories. Worth reading in full so you know exactly what was claimed, and what wasn’t.

Staying Safe

- Scamwatch, “How scammers use technology and AI,” <https://www.scamwatch.gov.au/stop-check-protect/help-to-spot-and-avoid-scams/how-scammers-use-technology-and-ai>, and “Why anyone can be a victim of a scam,” <https://www.scamwatch.gov.au/stop-check-protect/help-to-spot-and-avoid-scams/why-anyone-can-be-a-victim-of-a-scam>. The national scam service on voice cloning, deepfakes and fake messages, and why falling for a scam is not a sign of stupidity. It goes with Chapter 9.
- Australian Signals Directorate, “Social engineering,” <https://www.cyber.gov.au/threats/types-threats/social-engineering>. How criminals manipulate people rather than computers, and the habits that protect you. It goes with Chapter 9.
- eSafety Commissioner, “Image-based abuse,” <https://www.esafety.gov.au/key-topics/image-based-abuse>. What to do if intimate images of you, real or fake, are shared without consent, including how to report them. It goes with Chapter 9.
- eSafety Commissioner, *Talking to machines: Children’s experiences with AI assistants and companions* (2026), <https://www.esafety.gov.au/research/talking-to-machines-childrens-experiences-with-ai-assistants-and-companions>. Australian research on how children use AI chatbots and companions. Especially useful for parents and teachers. It goes with Chapters 6, 9, 13 and 14.

Creativity and Copyright

- Attorney-General’s Department, “Copyright and Artificial Intelligence Reference Group (CAIRG),” <https://www.ag.gov.au/rights-and-protections/copyright/copyright-and-artificial-intelligence-reference-group-cairg>. Where the Australian Government’s work on AI and copyright is published. The law here is still moving. It goes with Chapters 8 and 9.
- U.S. Copyright Office, *Copyright and Artificial Intelligence, Part 2: Copyrightability* (29 January 2025), <https://www.copyright.gov/ai/Copyright-and-Artificial-Intelligence-Part-2-Copyrightability-Report.pdf>. A readable official report on when work made with AI can be protected by copyright in the United States. It goes with Chapter 8.

Where AI Goes Next

- Department of Industry, Science and Resources, *National AI Plan* (2 December 2025), <https://www.industry.gov.au/publications/national-ai-plan>. The Australian Government’s plan for AI, including the Australian AI Safety Institute. It goes with Chapters 9, 11 and 14.

The Companion Website

- How-To-Use-AI.com, <https://how-to-use-ai.com>. Current tool walkthroughs, updated settings and screenshots, rechecked figures, and the parts of this book that change too quickly for print.

END OF PREVIEW

Thanks for reading this preview

You have been reading Chapters 1–3 of
AI for Normal People.
The full book continues with 11 more chapters.

Get the full book at
how-to-use-ai.com

Danijel-James Wynyard-McClay

How-To-Use-AI.com

AI for Normal People

You've already been using AI in your maps, your inbox, your bank and your streaming queue. This book explains what's actually going on, in plain English, without the hype or the jargon.

AI is software that finds patterns and makes predictions. Once you see it that way, the headlines get less alarming, the tools get more useful, and you can decide for yourself where AI belongs in your life and your work.

In this book you'll learn how to:

- Spot the AI you already use every day
- Understand what AI can and can't do, and why it sounds so sure of itself
- Ask AI tools for what you actually want
- Use AI at home and at work without handing over your judgement
- Tell real progress from hype



Danijel-James Wynyard-McClay

DJ is a software engineer, writer and podcaster based in Adelaide, Australia, and the creator of How To Use AI.com and Useful Stash. He builds software, automation and cloud infrastructure, likes evidence and strong opinions held loosely, and always asks “but does this actually fucking work?” before getting too excited.